

Monash Time-Series Forecasting Archive Replication

Fernando Urbano* Aben Carrington[†] Shrey Jain[‡]
Mukund Maheshwari[§]

March 11, 2024

Abstract

In this project we attempt to replicate results from a 2021 paper on the motivation and creation of the Monash Time Series Forecasting Archive, a project spearheaded by a group time series researchers from Monash University and the University of Sydney.

*fernandourbano@uchicago.edu

[†]acarrington@uchicago.edu

[‡]shreyjain@uchicago.edu

[§]mukundmaheshwari@uchicago.edu

1 Paper Importance

The Monash Archive is an essential asset to time series researchers as it provides a comprehensive benchmark time series forecasting archive to evaluate the performance of new global and multivariate forecasting algorithms. Specifically, as researchers branch further and further into the machine learning space, the Monash Archive allows them to test the generalized performance of their models against well-tested benchmark models, which is beneficial in addressing the questions of model overfitting and performance.

The archive contains datasets spanning multiple domains (industries) as well as 13 forecasting models, 6 of which are canonical univariate models, and 7 of which are global models that have shown positive results in recent years. In the following sections we provide a brief description of each of the datasets used, as well as an overview of the important aspects of the models used.

2 Data Sources

The archive repository contains 25 datasets where each data point is an individual time series, with most being of variable length. The datasets indicated as multivariate in Table 1 are constrained to have time series that are all of the same length so that multivariate forecasting models may be run on them without error. Additionally there are 5 datasets which contain a single very long time series. From these 30 primary datasets the authors have created 58 total datasets, where some are split according to differing model frequencies, and datasets with missing values are split into two datasets: one with and one without the missing values. It should be noted that global univariate and local univariate can be applied to all datasets in the repository.

We now go into further depth concerning each of the primary datasets:

2.1 Collections of multiple time series

This section describes the benchmark datasets that have a sufficient number of series from a particular frequency. The datasets may contain different categories in terms of domain and frequency.

2.1.1 M1 dataset

This dataset is from the M1 forecasting competition held in the year 1982. It contains 1001 time series with 3 different frequencies (monthly, quarterly, yearly) and is one of few belonging to multiple domains.

2.1.2 M3 dataset

This dataset is from the M3 forecasting competition held in the year 2000. It contains 3003 time series with 4 different frequencies (monthly, quarterly, yearly, other) and is one of few belonging to multiple domains.

2.1.3 M4 dataset

This dataset is from the M4 forecasting competition held in the year 2020. It contains 100,000 time series with 6 different frequencies (hourly, daily, weekly, monthly, quarterly, yearly) and is one of few belonging to multiple domains.

2.1.4 Tourism dataset

This dataset originates from a kaggle competition. It contains 1311 tourism related time series with 3 different frequencies (monthly, quarterly, yearly).

2.1.5 NN5 dataset

This dataset is from the NN5 neural forecasting competition held in the year 2008. It contains 111 daily time series of cash withdrawals from ATMs in the U.K. The original dataset contained missing values that were filled in by the authors of the paper using a median substitution method.

2.1.6 CIF 2016 dataset

This dataset is from the Computational Intelligence in Forecasting (CIF) competition held in 2016. It contains 72 monthly time series, where 24 originate from the banking sector and the other 48 are artificially generated. There are two datasets corresponding to different forecast horizons: 6 and 12.

2.1.7 Kaggle web traffic dataset

This dataset is from the Kaggle web traffic forecasting competition held in the year 2017. It contains 145063 daily time series representing the number of hits for a set of

pages on Wikipedia from 01/07/2015 to 10/09/2017. The authors also include their own aggregated version of weekly time series.

2.1.8 Solar dataset

This dataset corresponds to the solar power production in the state of Alabama throughout the year 2006. It contains 137 time series representing the amount of solar power produced every 10 minutes. The authors also include their own aggregated version of weekly time series.

2.1.9 Electricity dataset

This dataset corresponds to the amount of electricity consumed per hour by 321 clients, collected from 2012 to 2014 by UC Irvine. The authors also include their own aggregated version of weekly time series.

2.1.10 London smart meters dataset

This dataset corresponds to the energy consumption readings of London households in kWh from November 2011 to February 2014. It contains 5560 half-hourly time series. Two versions are included: one with missing values, and one where the missing values are filled in using the last observation carried forward (LOCF) method.

2.1.11 Australian electricity demand dataset

This dataset corresponds to the electricity demand of 5 Australian states: Victoria, New South Wales, Queensland, Tasmania and South Australia. It contains 5 half-hourly time series.

2.1.12 Wind farms dataset

This dataset contains very long minutely time series representing the wind power production of 339 wind farms in Australia. It is curated by the authors and is not available elsewhere. Two versions are included: one with missing values, and one where the missing values have been set to zero.

2.1.13 Car parts dataset

This dataset contains 2674 intermittent monthly time series showing car parts sales from January 1998 to March 2002. Two versions are included: one with missing

values, and one where the missing values have been set to zero.

2.1.14 Dominick dataset

This dataset corresponds to the profit of individual stock keeping units (SKUs) from a retailer collected from the online platform of the University of Chicago Booth School of Business Kilts Center. It contains 115704 weekly time series.

2.1.15 FRED-MD dataset

This dataset was extracted from the FRED-MD database and corresponds to a set of macro-economic indicators from the Federal Reserve Bank. It contains 107 monthly time series starting from 01/01/1959.

2.1.16 Bitcoin dataset

This dataset shows the potential factors influencing bitcoin price (such as transaction values and hash rate). It contains 18 daily time series, 2 of which show the public opinion of bitcoins in the form of tweets and google searches mentioning the keyword, bitcoin. It is curated by the authors and is not available elsewhere. Two versions are included: one with missing values, and one where the missing values are filled in using the LOCF method.

2.1.17 San Francisco Traffic dataset

This dataset corresponds to the road occupancy rates on San Francisco Bay area freeways. It contains 862 hourly time series taken from 2015 to 2016. The authors also include their own aggregated version of weekly time series.

2.1.18 Melbourne pedestrian counts dataset

This dataset contains hourly time series of pedestrian counts captured from 66 sensors in Melbourne from May 2009 to April 30, 2020.

2.1.19 Rideshare dataset

This dataset corresponds to attributes related to Uber and Lyft rideshare services (such as price and distance) for different locations in New York from 26/11/2018 to 18/12/2018. It contains 2304 hourly time series. Two versions are included: one with missing values, and one where the missing values have been set to zero.

2.1.20 Vehicle trips dataset

This dataset corresponds to the number of trips and vehicles belonging to a set of for-hire vehicle (FHV) companies in 2015, extracted from fivethirtyeight. It contains 329 daily time series. Two versions are included: one with missing values, and one where the missing values are filled in using the LOCF method.

2.1.21 Hospital dataset

This dataset corresponds to e patient counts related to medical products from January 2000 to December 2006. It contains 767 monthly time series.

2.1.22 COVID deaths dataset

This dataset represents the total COVID-19 deaths in a set of countries and states from 22/01/2020 to 20/08/2020, extracted from the Johns Hopkins repository. It contains 266 daily time series.

2.1.23 KDD cup 2018 dataset

This dataset originates from a 2018 competition. It contains 270 long hourly time series representing the air quality levels in 59 stations from 2 cities, Beijing (35 stations) and London (24 stations) from 01/01/2017 to 31/03/2018. It represents the air quality across multiple measurements.

2.1.24 Weather dataset

This dataset contains 3010 daily time series of four weather variables: rain, minimum temperature, maximum temperature, and solar radiation, measured at weather stations in Australia.

2.1.25 Temperature rain dataset

This dataset corresponds to the temperature/rainfall observations and forecasts, gathered by the Australian Bureau of Meteorology for 422 weather stations across Australia, between 02/05/2015 and 26/04/2017. It contains 32072 daily time series. Two versions are included: one with missing values, and one where the missing values have been set to zero.

2.2 Single long time series datasets

This section describes the benchmark datasets which have single time series with a large amount of data points.

2.2.1 Sunspot dataset

This dataset contains the single daily time series representing the sunspot numbers from 08/01/1818 to 31/05/2020. Two versions are included: one with missing values, and one where the missing values are filled in using the LOCF method.

2.2.2 Saugeen river flow dataset

This dataset contains a single very long time series representing the daily mean flow of the Saugeen River at Walkerton in cubic meters per second from 01/01/1915 to 31/12/1979. The length of the time series is 23,741.

2.2.3 US Births dataset

This dataset contains a single very long daily time series representing the number of births in the US from 01/01/1969 to 31/12/1988. The length of the time series is 7,305.

2.2.4 Solar power dataset

This dataset contains a single very long time series representing the solar power production of an Australian wind farm recorded every 4 seconds starting from 01/08/2019. The length of the time series is 7,397,222.

2.2.5 Wind power dataset

This dataset contains a single very long time series representing the wind power production of an Australian wind farm recorded every 4 seconds starting from 01/08/2019. The length of the time series is 7,397,147.

Table 1: Datasets in the current time series forecasting archive

	Dataset	Domain	No. of Series	Min. Length	Max. Length	No. of Freq	Missing	Competition	Multivariate
0	M1	Multiple	1023	18	150	3	No	Yes	No
1	M3	Multiple	3003	20	144	4	No	Yes	No
2	M4	Multiple	100000	19	9933	6	No	Yes	No
3	Tourism	Tourism	1311	11	333	3	No	Yes	No
4	CIF 2016	Banking	72	28	120	1	No	Yes	No
5	London Smart Meters	Energy	5560	288	39648	1	Yes	No	No
6	Aus. Electricity Demand	Energy	5	230736	232272	1	No	No	No
7	Wind Farms	Energy	339	6345	527040	1	Yes	No	No
8	Dominick	Sales	115704	28	393	1	No	No	No
9	Bitcoin	Economic	18	4581	4581	1	Yes	No	No
10	Pedestrian Counts	Transport	66	576	96424	1	No	No	No
11	Vehicle Trips	Transport	329	70	243	1	Yes	No	No
12	KDD Cup 2018	Nature	270	9504	10920	1	Yes	Yes	No
13	Weather	Nature	3010	1332	65981	1	No	No	No
14	NN5	Banking	111	791	791	2	Yes	Yes	Yes
15	Web Traffic	Web	145063	803	803	2	Yes	Yes	Yes
16	Solar	Energy	137	52560	52560	2	No	No	Yes
17	Electricity	Energy	321	26304	26304	2	No	No	Yes
18	Car Parts	Sales	2674	51	51	1	Yes	No	Yes
19	FRED-MD	Economic	107	728	728	1	No	No	Yes
20	San Francisco Traffic	Transport	862	17544	17544	2	No	No	Yes
21	Rideshare	Transport	2304	541	541	1	Yes	No	Yes
22	Hospital	Health	767	84	84	1	No	No	Yes
23	COVID Deaths	Nature	266	212	212	1	No	No	Yes
24	Temperature Rain	Nature	32072	725	725	1	Yes	No	Yes
25	Sunspot	Nature	1	73924	73924	1	Yes	No	No
26	Saugeen River Flow	Nature	1	23741	23741	1	No	No	No
27	US Births	Nature	1	7305	7305	1	No	No	No
28	Solar Power	Energy	1	7397222	7397222	1	No	No	No
29	Wind Power	Energy	1	7397147	7397147	1	No	No	No

3 Models/Evaluation

3.1 Models

As previously mentioned this project uses 6 traditional univariate models, and 7 global models, covering a representative set of state-of-the-art forecasting models from statistical, machine learning, and deep learning domains, for a total of 13 models.

The 6 traditional models used are Exponential Smoothing (ETS), Auto-Regressive Integrated Moving Average (ARIMA), Simple Exponential Smoothing (SES), Theta, Trigonometric Box-Cox ARMA Trend Seasonal (TBATS), and Dynamic Harmonic Regression ARIMA (DHR-ARIMA). The 7 global forecasting models used are a linear Pooled Regression model (PR), a Feed-Forward Neural Network (FFNN), CatBoost, DeepAR, N-BEATS, a WaveNet, and a Transformer method. As there is extensive literature on each of these models we forego any description here.

We implement the 6 traditional univariate models as well as PR and CatBoost

in R using the packages forecast, glmnet, and catboost. The authors of the original paper used R base version 4.0.2 but we use R base version 4.3.2. The other models are implemented in Python using the GluonTS package from AWS. The authors of the original paper used Python 3.7.4 and GluonTS 0.8.0. This presented challenges for us that will be expounded upon later. Since all models are presented as benchmarks for baseline model performance, no hyperparameter tuning is done and the models are presented with their default hyperparameters.

3.2 Model Evaluation

For evaluating the performance of the models, the authors compared the Mean Absolute Scaled Error (MASE) of each of the models per dataset. This statistic was calculated by using forecasting functions created by the authors of the original paper (and publicly available on github) to calculate the MASE per time series in each of the datasets and then calculating the mean value across the MASE results per dataset. The formula for this is given below

$$\text{MASE} = \frac{\sum_{k=M+1}^{M+h} |F_k - Y_k|}{\frac{h}{M-S} \sum_{k=S+1}^M |Y_{k+1} - Y_k|}$$

where M is the number of data points in the training series, S is the seasonality of the dataset, h is the forecast horizon, F_k are the generated forecasts, and Y_k are the actual values.

The MASE results per model (columns) per dataset (rows) are given in Table 2.

Table 2: Mean MASE results. The best model across each dataset is highlighted in boldface.

	Dataset	SES	Theta	ETS	(DHR-) ARIMA	PR	Cat Boost	ARIMA	TBATS
0	M1 Yearly	4.938	4.191	3.771	-	4.588	4.333	4.479	3.499
1	M1 Quarterly	1.929	1.702	1.658	-	1.892	2.040	1.787	1.694
2	M1 Monthly	1.379	1.091	1.074	-	1.123	1.220	1.165	1.118
3	M3 Quarterly	1.417	1.117	1.170	-	1.248	1.449	1.240	1.256
4	M3 Monthly	1.091	0.864	0.865	-	1.010	1.076	0.873	0.861
5	M4 Yearly	-	-	-	-	-	-	-	3.437
6	M4 Quarterly	-	-	-	-	-	-	-	1.186
7	M4 Weekly	-	-	-	-	-	-	-	0.505
8	Tourism Yearly	3.253	3.015	3.395	-	3.516	3.619	3.775	3.685
9	Tourism Quarterly	3.210	1.661	1.592	-	1.643	1.821	1.776	1.833
10	Tourism Monthly	3.306	1.649	1.526	-	1.678	1.712	1.587	1.751
11	Bitcoin	5.289	5.223	4.538	-	4.616	5.653	5.498	4.602
12	Vehicle Trips	2.273	1.914	1.964	-	2.196	2.004	2.051	1.856
13	NN5 Daily	-	-	-	-	-	0.970	-	0.858
14	NN5 Weekly	0.903	0.885	0.911	0.887	0.854	0.854	-	0.872
15	Solar Weekly	1.215	1.224	1.134	0.848	1.053	1.477	-	0.916
16	Electricity Weekly	1.536	1.476	1.526	0.878	0.916	0.813	-	0.792
17	Traffic Weekly	1.116	1.121	1.125	1.191	1.122	1.122	-	1.148
18	Hospital	0.813	0.761	0.765	-	0.782	0.796	0.788	0.768
19	Sunspot	0.128	0.128	0.128	-	0.099	0.073	0.067	0.064
20	CIF 2016	1.291	0.997	0.841	-	1.019	1.200	0.927	0.861
21	COVID Deaths	7.776	7.793	5.326	-	8.731	8.092	6.104	5.719
22	Car Parts	0.897	0.914	0.925	-	0.755	0.853	0.927	1.002
23	Fred Md	0.617	0.698	0.468	-	8.827	0.988	0.532	0.502
24	M3 Yearly	3.167	2.774	2.860	-	3.223	3.711	3.417	3.127
25	Saugeen River Flow	1.426	1.425	2.036	-	1.674	1.430	1.548	1.477
26	US Births	4.343	2.138	1.529	-	2.094	1.690	1.917	1.453

As extra material, we also create the Tables with median and mean MAE, RMSE, sMAPE. Similarly to MASE, sMAPE normalizes errors by the sum of the actual and predicted values, making it scale-invariant in terms of the magnitude of the data. On the other hand, for RMSE and sMAPE, the results are not directly comparable between datasets, but only within each dataset.

Table 3: Mean MAE results. The best model across each dataset is highlighted in boldface.

	Dataset	SES	Theta	ETS	(DHR-) ARIMA	PR	Cat Boost	ARIMA	TBATS
0	M1 Yearly	171.4K	152.8K	146.1K	-	134.2K	249.4K	145.6K	103.0K
1	M1 Quarterly	2.2K	2.0K	2.1K	-	1.6K	1.9K	2.2K	2.3K
2	M1 Monthly	2.3K	2.2K	1.9K	-	2.1K	2.1K	2.1K	2.2K
3	M3 Quarterly	572	486	513	-	519	594	559	562
4	M3 Monthly	743	624	626	-	693	736	655	631
5	M4 Yearly	-	-	-	-	-	-	-	960
6	M4 Quarterly	-	-	-	-	-	-	-	570
7	M4 Weekly	-	-	-	-	-	-	-	297
8	Tourism Yearly	95.6K	90.7K	94.8K	-	82.7K	81.3K	95.0K	94.1K
9	Tourism Quarterly	15.0K	7.7K	8.9K	-	9.1K	10.1K	10.4K	10.0K
10	Tourism Monthly	5.3K	2.1K	2.0K	-	2.2K	2.5K	2.5K	2.9K
11	Bitcoin	1773.4Qi	1773.4Qi	1103.6Qi	-	666.4Qa	1921.1Qi	1047.2Qi	990.4Qa
12	Vehicle Trips	29.980	23.299	21.258	-	27.243	22.732	23.456	21.045
13	NN5 Daily	-	-	-	-	-	4.200	-	3.701
14	NN5 Weekly	15.665	15.305	15.698	15.383	14.937	15.359	-	14.985
15	Solar Weekly	1.2K	1.2K	1.1K	840	1.0K	1.5K	-	909
16	Electricity Weekly	74.1K	74.1K	67.7K	28.5K	44.9K	34.7K	-	24.4K
17	Traffic Weekly	1.125	1.131	1.144	1.222	1.125	1.181	-	1.166
18	Hospital	21.761	18.539	17.966	-	19.237	19.114	19.742	17.429
19	Sunspot	4.933	4.933	4.933	-	3.833	2.800	2.567	2.467
20	CIF 2016	581.9K	714.8K	642.4K	-	563.2K	688.0K	469.1K	855.6K
21	COVID Deaths	354	321	85.591	-	348	486	85.768	96.288
22	Car Parts	0.548	0.530	0.564	-	0.407	0.531	0.561	0.583
23	Fred Md	2.8K	3.5K	2.0K	-	8.9K	2.6K	3.0K	2.0K
24	M3 Yearly	1.0K	957	1.0K	-	1.0K	1.1K	1.4K	1.2K
25	Saugeen River Flow	21.497	21.486	30.693	-	25.241	21.562	23.338	22.262
26	US Births	1.2K	587	420	-	575	464	526	399

Table 4: Mean RMSE results. The best model across each dataset is highlighted in boldface.

	Dataset	SES	Theta	ETS	(DHR-) ARIMA	PR	Cat Boost	ARIMA	TBATS
0	M1 Yearly	193.8K	171.5K	167.7K	-	152.0K	269.0K	175.3K	116.9K
1	M1 Quarterly	2.5K	2.3K	2.4K	-	1.9K	2.2K	2.5K	2.7K
2	M1 Monthly	2.7K	2.6K	2.3K	-	2.5K	2.5K	2.5K	2.6K
3	M3 Quarterly	671	568	599	-	606	698	651	654
4	M3 Monthly	894	754	755	-	830	879	791	765
5	M4 Yearly	-	-	-	-	-	-	-	1.1K
6	M4 Quarterly	-	-	-	-	-	-	-	673
7	M4 Weekly	-	-	-	-	-	-	-	358
8	Tourism Yearly	106.7K	99.9K	104.7K	-	89.6K	89.6K	106.1K	105.8K
9	Tourism Quarterly	17.3K	9.3K	10.8K	-	11.7K	12.6K	12.5K	12.0K
10	Tourism Monthly	7.0K	2.7K	2.5K	-	2.7K	3.1K	3.1K	3.7K
11	Bitcoin	1963.7Qi	1963.7Qi	1223.5Qi	-	829.2Qa	2002.0Qi	1198.1Qi	1164.3Qi
12	Vehicle Trips	36.525	27.814	26.153	-	31.692	27.348	28.535	25.503
13	NN5 Daily	-	-	-	-	-	5.715	-	5.204
14	NN5 Weekly	18.825	18.647	18.816	18.550	18.615	18.711	-	18.528
15	Solar Weekly	1.3K	1.3K	1.3K	968	1.2K	1.7K	-	1.0K
16	Electricity Weekly	77.1K	76.9K	70.4K	32.6K	47.8K	37.6K	-	28.0K
17	Traffic Weekly	1.514	1.529	1.534	1.545	1.503	1.511	-	1.528
18	Hospital	26.551	22.592	22.023	-	23.479	23.287	23.837	21.281
19	Sunspot	4.946	4.946	4.946	-	3.954	3.141	2.938	2.595
20	CIF 2016	657.1K	804.7K	722.4K	-	648.9K	760.0K	526.4K	940.1K
21	COVID Deaths	403	370	102	-	394	617	100	113
22	Car Parts	0.784	0.782	0.802	-	0.729	0.794	0.811	0.837
23	Fred Md	3.1K	3.9K	2.3K	-	9.7K	2.8K	3.3K	2.3K
24	M3 Yearly	1.2K	1.1K	1.2K	-	1.2K	1.3K	1.7K	1.4K
25	Saugeen River Flow	39.794	39.787	50.392	-	47.703	39.306	45.536	42.576
26	US Births	1.4K	736	607	-	732	635	706	607

Table 5: Mean sMAPE results. The best model across each dataset is highlighted in boldface.

	Dataset	SES	Theta	ETS	(DHR-) ARIMA	PR	Cat Boost	ARIMA	TBATS
0	M1 Yearly	0.231	0.202	0.186	-	0.188	0.200	0.195	0.174
1	M1 Quarterly	0.181	0.163	0.174	-	0.166	0.177	0.166	0.166
2	M1 Monthly	0.171	0.155	0.146	-	0.148	0.162	0.153	0.148
3	M3 Quarterly	0.109	0.092	0.097	-	0.098	0.112	0.102	0.102
4	M3 Monthly	0.162	0.139	0.141	-	0.152	0.165	0.143	0.138
5	M4 Yearly	-	-	-	-	-	-	-	0.149
6	M4 Quarterly	-	-	-	-	-	-	-	0.102
7	M4 Weekly	-	-	-	-	-	-	-	0.073
8	Tourism Yearly	0.341	0.319	0.365	-	0.469	0.328	0.334	0.339
9	Tourism Quarterly	0.274	0.154	0.151	-	0.159	0.167	0.165	0.172
10	Tourism Monthly	0.364	0.199	0.190	-	0.211	0.213	0.196	0.212
11	Bitcoin	0.208	0.302	0.191	-	0.215	0.306	0.269	0.200
12	Vehicle Trips	0.362	0.301	0.313	-	0.350	0.308	0.308	0.291
13	NN5 Daily	-	-	-	-	-	0.239	-	0.211
14	NN5 Weekly	0.122	0.120	0.123	0.118	0.114	0.117	-	0.116
15	Solar Weekly	0.246	0.248	0.229	0.179	0.217	0.285	-	0.191
16	Electricity Weekly	0.142	0.146	0.141	0.108	0.100	0.097	-	0.085
17	Traffic Weekly	0.124	0.125	0.126	0.134	0.125	0.130	-	0.128
18	Hospital	0.179	0.173	0.175	-	0.176	0.179	0.178	0.176
19	Sunspot	1.924	1.924	1.924	-	1.901	1.858	1.730	1.860
20	CIF 2016	0.149	0.130	0.122	-	0.123	0.151	0.114	0.122
21	COVID Deaths	0.153	0.156	0.086	-	0.183	0.158	0.092	0.087
22	Car Parts	0.649	0.593	0.658	-	0.432	0.655	0.657	0.659
23	Fred Md	0.087	0.097	0.084	-	0.308	0.093	0.080	0.080
24	M3 Yearly	0.178	0.168	0.170	-	0.171	0.197	0.188	0.174
25	Saugeen River Flow	0.360	0.360	0.675	-	0.453	0.362	0.398	0.373
26	US Births	0.118	0.058	0.041	-	0.058	0.045	0.052	0.038

Table 6: Median MAE results. The best model across each dataset is highlighted in boldface.

	Dataset	SES	Theta	ETS	(DHR-) ARIMA	PR	Cat Boost	ARIMA	TBATS
0	M1 Yearly	379	256	191	-	246	261	180	173
1	M1 Quarterly	22.296	19.554	19.588	-	19.195	19.804	16.228	18.871
2	M1 Monthly	45.333	38.230	38.508	-	37.365	39.924	40.538	35.776
3	M3 Quarterly	372	294	305	-	325	397	334	336
4	M3 Monthly	517	421	409	-	479	533	412	407
5	M4 Yearly	-	-	-	-	-	-	-	430
6	M4 Quarterly	-	-	-	-	-	-	-	256
7	M4 Weekly	-	-	-	-	-	-	-	164
8	Tourism Yearly	4.3K	4.1K	4.3K	-	4.3K	5.0K	4.6K	4.8K
9	Tourism Quarterly	1.9K	1.1K	1.0K	-	992	1.0K	1.0K	1.2K
10	Tourism Monthly	968	478	457	-	475	473	463	492
11	Bitcoin	23.2K	20.3K	19.4K	-	25.1K	20.9K	29.9K	27.3K
12	Vehicle Trips	6.033	4.667	4.667	-	6.967	5.367	4.967	4.433
13	NN5 Daily	-	-	-	-	-	3.684	-	3.458
14	NN5 Weekly	14.183	13.904	14.273	14.824	12.837	13.129	-	13.727
15	Solar Weekly	1.1K	1.1K	1.1K	761	942	1.3K	-	780
16	Electricity Weekly	11.0K	10.4K	11.0K	6.8K	7.1K	6.1K	-	6.1K
17	Traffic Weekly	0.918	0.924	0.918	0.976	0.930	0.948	-	0.942
18	Hospital	6.667	6.667	6.667	-	6.667	6.917	6.833	6.833
19	Sunspot	4.933	4.933	4.933	-	3.833	2.800	2.567	2.467
20	CIF 2016	107	103	70.431	-	95.132	111	80.656	67.118
21	COVID Deaths	2.233	4.417	1.650	-	6.767	3.217	1.783	1.800
22	Car Parts	0.333	0.250	0.333	-	0.250	0.417	0.333	0.417
23	Fred Md	1.894	1.940	2.350	-	41.359	4.114	2.732	1.992
24	M3 Yearly	703	660	641	-	712	860	701	638
25	Saugeen River Flow	21.497	21.486	30.693	-	25.241	21.562	23.338	22.262
26	US Births	1.2K	587	420	-	575	464	526	399

Table 7: Median MASE results. The best model across each dataset is highlighted in boldface.

	Dataset	SES	Theta	ETS	(DHR-) ARIMA	PR	Cat Boost	ARIMA	TBATS
0	M1 Yearly	3.772	3.155	2.324	-	2.847	2.912	2.127	2.215
1	M1 Quarterly	1.417	1.264	1.196	-	1.376	1.411	1.171	1.200
2	M1 Monthly	1.167	0.885	0.851	-	0.947	1.016	0.896	0.902
3	M3 Quarterly	1.073	0.831	0.855	-	0.902	1.126	0.917	0.914
4	M3 Monthly	0.861	0.721	0.712	-	0.825	0.900	0.704	0.699
5	M4 Yearly	-	-	-	-	-	-	-	2.402
6	M4 Quarterly	-	-	-	-	-	-	-	0.915
7	M4 Weekly	-	-	-	-	-	-	-	0.365
8	Tourism Yearly	2.442	2.360	2.373	-	2.356	3.000	2.719	2.518
9	Tourism Quarterly	2.309	1.348	1.275	-	1.361	1.368	1.388	1.477
10	Tourism Monthly	2.336	1.382	1.276	-	1.484	1.461	1.333	1.491
11	Bitcoin	3.089	2.955	2.686	-	3.166	3.018	3.542	3.207
12	Vehicle Trips	1.402	0.999	0.964	-	1.429	1.129	1.020	0.963
13	NN5 Daily	-	-	-	-	-	0.902	-	0.834
14	NN5 Weekly	0.781	0.805	0.775	0.769	0.781	0.808	-	0.827
15	Solar Weekly	1.231	1.241	1.209	0.861	1.063	1.475	-	0.894
16	Electricity Weekly	1.341	1.303	1.337	0.798	0.842	0.732	-	0.705
17	Traffic Weekly	0.973	0.983	0.977	1.035	0.980	0.946	-	0.996
18	Hospital	0.745	0.723	0.731	-	0.740	0.754	0.736	0.734
19	Sunspot	0.128	0.128	0.128	-	0.099	0.073	0.067	0.064
20	CIF 2016	0.862	0.662	0.532	-	0.746	0.861	0.559	0.537
21	COVID Deaths	1.554	2.192	0.614	-	5.313	2.052	0.982	0.605
22	Car Parts	0.562	0.482	0.562	-	0.375	0.562	0.600	0.596
23	Fred Md	0.430	0.407	0.385	-	8.458	0.618	0.355	0.370
24	M3 Yearly	2.261	1.985	1.907	-	2.267	2.726	2.003	1.900
25	Saugeen River Flow	1.426	1.425	2.036	-	1.674	1.430	1.548	1.477
26	US Births	4.343	2.138	1.529	-	2.094	1.690	1.917	1.453

Table 8: Median RMSE results. The best model across each dataset is highlighted in boldface.

	Dataset	SES	Theta	ETS	(DHR-) ARIMA	PR	Cat Boost	ARIMA	TBATS
0	M1 Yearly	416	323	230	-	305	298	208	204
1	M1 Quarterly	24.459	22.811	21.858	-	22.529	22.572	20.232	22.320
2	M1 Monthly	54.669	46.396	44.392	-	45.346	47.584	47.105	44.038
3	M3 Quarterly	436	356	369	-	378	480	406	400
4	M3 Monthly	634	517	496	-	582	634	500	493
5	M4 Yearly	-	-	-	-	-	-	-	495
6	M4 Quarterly	-	-	-	-	-	-	-	302
7	M4 Weekly	-	-	-	-	-	-	-	200
8	Tourism Yearly	4.7K	4.6K	4.6K	-	4.7K	5.5K	5.2K	5.2K
9	Tourism Quarterly	2.3K	1.4K	1.2K	-	1.2K	1.2K	1.2K	1.5K
10	Tourism Monthly	1.3K	675	599	-	596	620	606	671
11	Bitcoin	30.3K	26.3K	24.3K	-	31.4K	28.3K	38.2K	33.0K
12	Vehicle Trips	8.103	5.802	5.925	-	8.725	6.962	6.506	5.580
13	NN5 Daily	-	-	-	-	-	5.320	-	4.749
14	NN5 Weekly	17.524	16.816	17.523	17.487	16.263	16.060	-	16.990
15	Solar Weekly	1.2K	1.2K	1.2K	878	1.0K	1.4K	-	886
16	Electricity Weekly	12.5K	11.8K	12.5K	8.3K	8.2K	7.3K	-	7.3K
17	Traffic Weekly	1.201	1.215	1.210	1.211	1.195	1.159	-	1.214
18	Hospital	8.256	8.196	8.251	-	8.251	8.485	8.391	8.357
19	Sunspot	4.946	4.946	4.946	-	3.954	3.141	2.938	2.595
20	CIF 2016	129	118	85.771	-	109	131	103	79.025
21	COVID Deaths	3.087	5.290	2.205	-	8.283	3.941	2.164	2.129
22	Car Parts	0.707	0.645	0.707	-	0.577	0.707	0.707	0.707
23	Fred Md	2.306	2.362	2.702	-	45.182	4.512	3.490	2.515
24	M3 Yearly	804	740	759	-	825	969	815	753
25	Saugeen River Flow	39.794	39.787	50.392	-	47.703	39.306	45.536	42.576
26	US Births	1.4K	736	607	-	732	635	706	607

Table 9: Median sMAPE results. The best model across each dataset is highlighted in boldface.

	Dataset	SES	Theta	ETS	(DHR-) ARIMA	PR	Cat Boost	ARIMA	TBATS
0	M1 Yearly	0.173	0.147	0.130	-	0.135	0.134	0.120	0.127
1	M1 Quarterly	0.112	0.086	0.084	-	0.101	0.116	0.097	0.086
2	M1 Monthly	0.143	0.112	0.108	-	0.119	0.125	0.115	0.113
3	M3 Quarterly	0.067	0.052	0.055	-	0.057	0.076	0.064	0.062
4	M3 Monthly	0.107	0.093	0.091	-	0.104	0.110	0.090	0.090
5	M4 Yearly	-	-	-	-	-	-	-	0.088
6	M4 Quarterly	-	-	-	-	-	-	-	0.058
7	M4 Weekly	-	-	-	-	-	-	-	0.048
8	Tourism Yearly	0.188	0.168	0.192	-	0.169	0.236	0.227	0.206
9	Tourism Quarterly	0.225	0.132	0.129	-	0.133	0.135	0.131	0.148
10	Tourism Monthly	0.302	0.174	0.172	-	0.185	0.189	0.180	0.190
11	Bitcoin	0.182	0.187	0.188	-	0.172	0.187	0.192	0.175
12	Vehicle Trips	0.342	0.235	0.232	-	0.327	0.271	0.236	0.228
13	NN5 Daily	-	-	-	-	-	0.229	-	0.196
14	NN5 Weekly	0.109	0.110	0.108	0.111	0.105	0.104	-	0.110
15	Solar Weekly	0.248	0.249	0.244	0.176	0.218	0.282	-	0.184
16	Electricity Weekly	-	0.117	-	0.070	-	0.061	-	-
17	Traffic Weekly	0.097	0.098	0.098	0.105	0.098	0.102	-	0.101
18	Hospital	0.166	0.159	0.161	-	0.161	0.168	0.168	0.163
19	Sunspot	1.962	1.962	1.962	-	1.956	1.933	1.943	1.947
20	CIF 2016	0.114	0.080	0.066	-	0.084	0.108	0.077	0.070
21	COVID Deaths	-	-	-	-	-	-	-	-
22	Car Parts	-	-	-	-	-	-	-	-
23	Fred Md	0.016	0.015	0.015	-	0.291	0.033	0.016	0.013
24	M3 Yearly	0.124	0.115	0.115	-	0.129	0.146	0.124	0.115
25	Saugeen River Flow	0.360	0.360	0.676	-	0.454	0.363	0.398	0.374
26	US Births	0.118	0.058	0.041	-	0.058	0.045	0.052	0.038

4 Replication Performance

The goal of our replication project was to successfully load the datasets, implement the authors functions with minimal updates, and regenerate Table 1 and Table 2 from the original paper in similar formatting. This proved to be a project primarily rooted in understanding effective package managing and resolving software conflicts related to cross-package version dependencies. We were successfully able to run all the models using R on a significant portion of the datasets but had lingering, unresolvable issues on the GluonTS models and some datasets. We will now briefly describe these successes, failures, and some recommendations, given our current knowledge, for the future of this replication project.

4.1 Successes

We successfully generated Table 1, having minor differences for 3 cells in comparison with the original model due to (i) Addition of new datasets, (ii) addition of new data inside the time-series.

In Table 2, we were successfully able to run the R models on a significant portion of the datasets.

We used an Anaconda environment and initially encountered problems with package installations from the authors’ indicated R package versions resulting from the fact that the authors originally used an R base version of 4.0.2, which was unable to be installed in current versions of Anaconda. We instead used an R base version of 4.3.2 and were able to bypass these restrictions by manually adding the packages in our conda environment through including several `"install.packages()"` statements in our R forecasting scripts. However, there was an issue in building one of the dependency packages resulting from system incompatibilities with the underlying C and Fortran used to write the packages. We found that in order to successfully build the packages the user must first install cmake through homebrew on their machine. We also found through resolving these issues that it is essential to use the libmamba solver in one’s conda environment as opposed to the classic solver.

Additionally, when we were finally able to run the R models on the datasets, we found that some models took an infeasible amount of time to run on local machines for certain datasets (on the order of 4 hours for a single model on a single dataset). This indicated to us that it would be wiser to run these models on an HPC cluster. However, as non faculty or staff researchers, the University of Chicago will not grant us access to the Midway Clusters and so we were unable to run these datasets in the final product in the interest of time and reproducibility.

For the datasets that we were successfully able to run, we found equal results for over 95% of the models (perfect precision)!

4.2 Failures

We were not as fortunate in running the Python models as we were in running the R models. In addition to our inability to run a few of the datasets, we had several issues related to Python package versioning primarily resulting from Conda’s poor ability to run versions of Python and Python packages that are more than 2 years out-of-date (currently, the Python packages the paper’s authors used are nearly 3 years out-of-date).

These issues began with a massive update in the GluonTS package between 2021 (version 0.8.0) and today (version 0.14.4) which significantly altered the declarations

of model objects. This update made it so that the authors' model functions were declaring functional arguments for the GluonTS model objects that no longer existed in the GluonTS documentation. As GluonTS is primarily based on the complex machine learning libraries PyTorch and MXnet, and none of our group members had the knowledge of these libraries to update the authors' functions, our only option was to attempt to downgrade the version of GluonTS from 0.14.4 to 0.8.0 and resolve any resulting dependency conflicts in the Conda environment.

Our group member attempting to resolve these issues was working using a homebrew installation of the miniforge distribution of mamba, while the rest of the group was working in the traditional installation of Anaconda3. The miniforge distribution provided much more versatility, however we required, based on the parameters of this project, that any solution be implementable using the conventional installation of Anaconda3. We were able to find a package configuration that passed running "pip check" for dependency conflicts and contained package versions very close to the package requirements of GluonTS version 0.8.0 (which was found on github). However, this configuration required the version of Python to be 3.7.6 (close to the author's version of Python 3.7.4), and the oldest version of Python available given the new anaconda update to Conda version 24.1.2 (which is required to run the R models), was Python 3.8.5. This presented an insurmountable conflict that we were not able to overcome without overhauling the entire project and starting from scratch using an alternate package manager.

4.3 Recommendations

For attempting this project again in the future, we recommend using either the miniforge distribution of mamba, or an entirely different package manager from Anaconda altogether. In researching our dependency conflicts, we found that the main advantage of the Miniconda distribution as opposed to the full Anaconda distribution is that Miniconda only ships with the repository management system as well as a limited number of base python packages, whereas Anaconda ships with 150+ python packages as well as several other modules (AWS, JupyterLab, JupyterNotebook, etc.) that are unnecessary for this project. Thus, the limited available versions of python and necessary python packages in Anaconda results from needless incompatibilities with system packages that aren't even used in this project. In light of this information, the best course of action may be to simply use pip as the package manager in a local virtual environment which bypasses Conda altogether, as this would provide the highest degree of versatility in using outdated package versions.